

Case-Indexer Service Technical Report (CORRECTED)

Executive Summary

The Case-Indexer is a **Kafka-to-gRPC bridge service** in Datadog's case management domain that serves as an asynchronous message processor. It consumes events from the `case-intake` Kafka topic and routes them to appropriate backend services (case-api and timeline-api) via gRPC calls. **The service does NOT directly integrate with Elasticsearch** - all database and search operations are handled by the downstream gRPC services.

Service Overview

Primary Functions

1. **Kafka Consumer:** Processes messages from `case-intake` topic with 32 partitions
2. **Message Router:** Routes different request types to appropriate gRPC services
3. **Async Bridge:** Provides asynchronous processing for case management operations
4. **Error Handler:** Implements retry logic with exponential backoff and pod restart strategies

Key Metrics

- **Topic:** `case-intake` with 32 partitions
- **Consumer Group:** `case-indexer`
- **Deployment:** 2 replicas per datacenter
- **Processing:** Up to 32 parallel workers (one per partition)
- **Request Types:** 4 main types routed via gRPC

Architecture Overview

System Architecture

The case-indexer follows a simple **message consumer and router pattern**: - Data Sources publish to Kafka topic - **Case-Indexer** consumes messages and makes gRPC calls - **Backend Services** (case-api, timeline-api) handle all data persistence and Elasticsearch operations - **case-indexer has NO direct database or Elasticsearch connections**

Data Flow Architecture

IMPORTANT: The case-indexer is purely a **Kafka→gRPC router**: 1. **Client Applications** publish requests to `case-intake` topic 2. **case-indexer** pulls messages and decodes protobuf payloads 3. **gRPC calls** to case-api or timeline-api services 4. **Backend services** handle PostgreSQL storage and Elasticsearch indexing 5. **case-indexer** only handles acknowledgment and error retry logic

Event Processing Pipeline

Message Processing Flow

The processing pipeline is straightforward: 1. **Message Decoding:** Decode intake payload from Kafka 2. **Request Type Detection:** Determine which gRPC service to call 3. **gRPC Service Call:** Route to case-api or timeline-api 4. **Error Classification:** Handle gRPC response errors 5. **Retry Logic:** Exponential backoff or pod restart for failures

Request Types and gRPC Routing

1. CreateRequest (*casepb.CreateRequest)

- **gRPC Call:** `caseClient.Create(ctx, req)`
- **Backend:** case-api service handles PostgreSQL and Elasticsearch operations
- **case-indexer role:** Message routing only

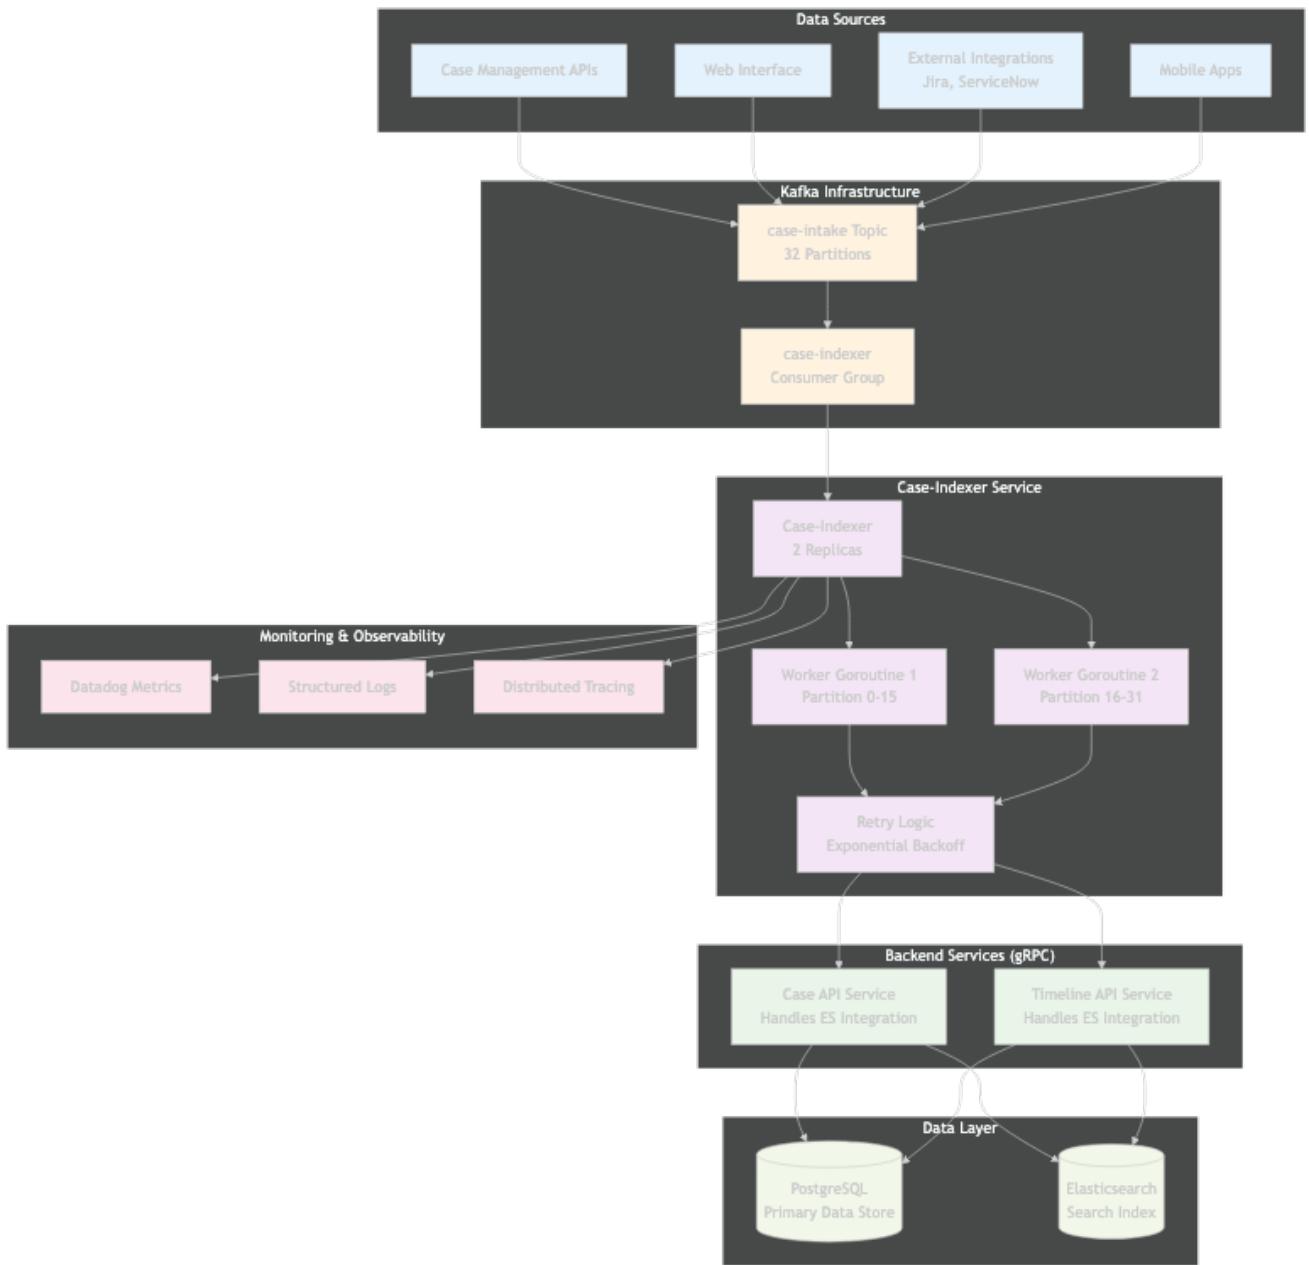


Figure 1: System Architecture

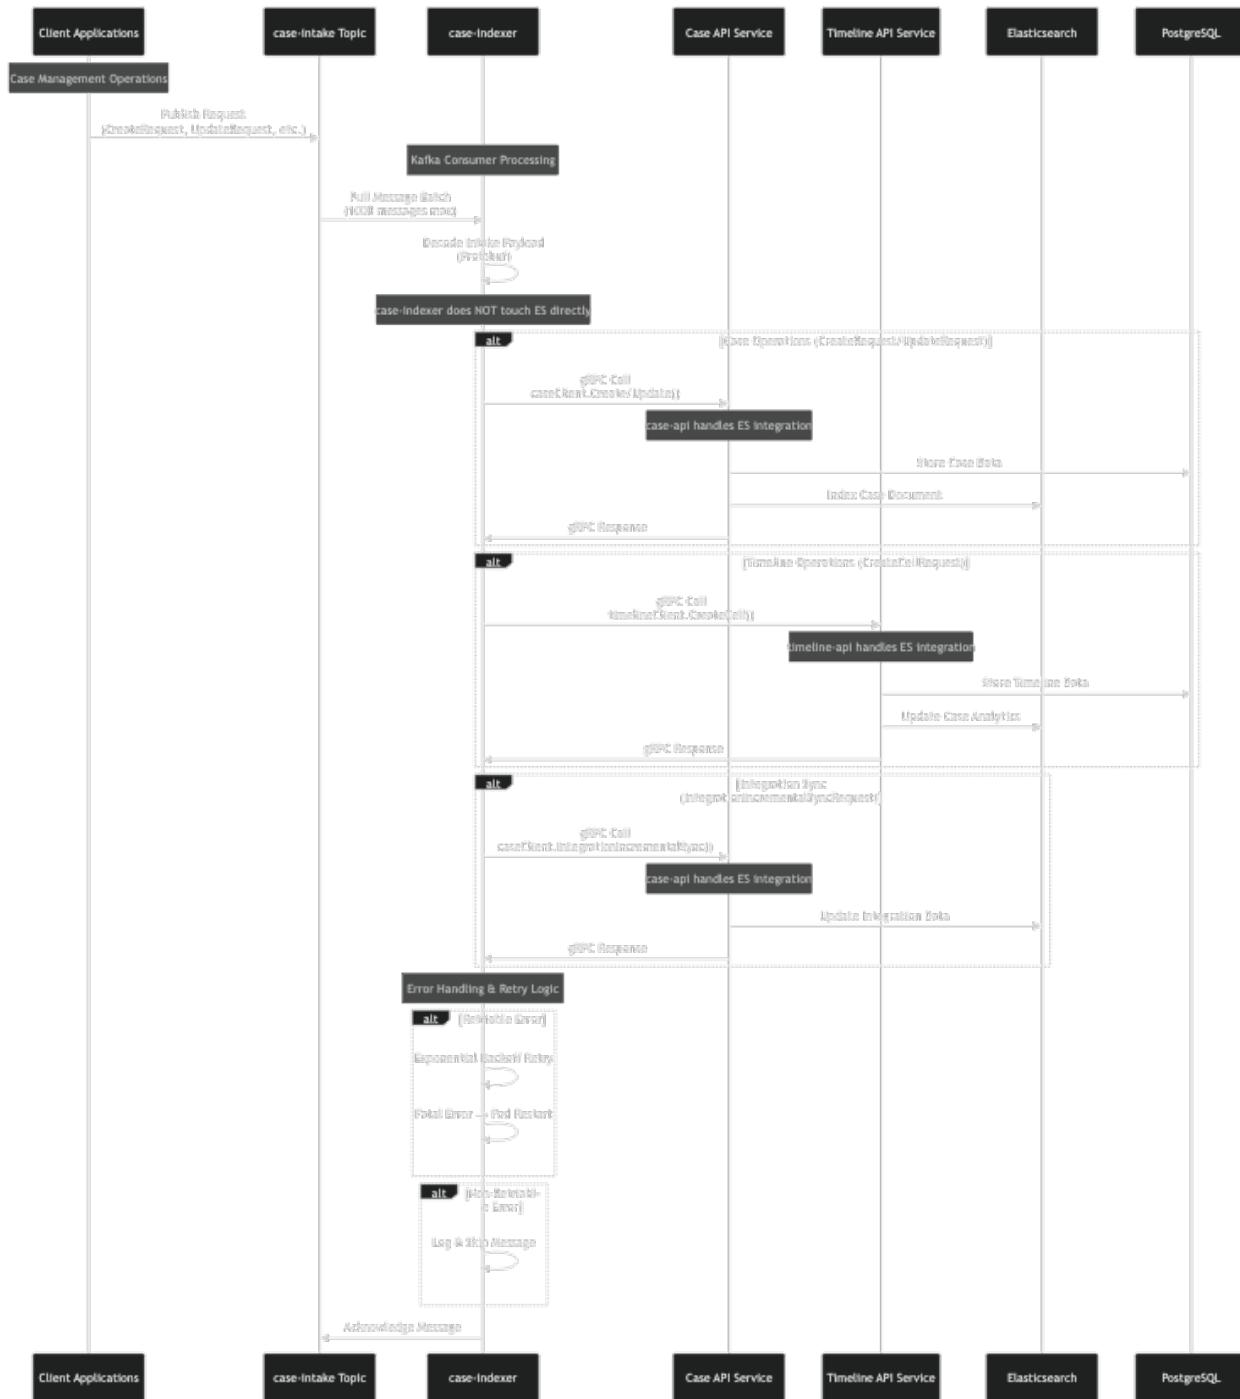


Figure 2: Data Flow

Event Processing Pipeline

Figure 3: Event Processing Pipeline

2. **UpdateRequest** (`*casepb.UpdateRequest`)
 - **gRPC Call:** `caseClient.Update(ctx, req)`
 - **Backend:** case-api service handles data updates and reindexing
 - **case-indexer role:** Message routing only

3. **CreateCellRequest** (`*timelinepb.CreateCellRequest`)
 - **gRPC Call:** `timelineClient.CreateCell(ctx, req)`
 - **Backend:** timeline-api service handles timeline storage and analytics updates
 - **case-indexer role:** Message routing only

4. **IntegrationIncrementalSyncRequest** (`*casepb.IntegrationIncrementalSyncRequest`)
 - **gRPC Call:** `caseClient.IntegrationIncrementalSync(ctx, req)`
 - **Backend:** case-api service handles integration data synchronization
 - **case-indexer role:** Message routing only

What case-indexer Does NOT Do

No Direct Database Integration

- **No PostgreSQL connections:** All database operations via gRPC
- **No Elasticsearch client:** All search operations via case-api/timeline-api
- **No direct data manipulation:** Pure message routing service

No Business Logic

- **No data transformation:** Passes protobuf messages as-is to gRPC services
- **No validation:** Relies on backend services for request validation
- **No caching:** Stateless message processor

No Direct External Integrations

- **No Jira/ServiceNow clients:** Integration logic in backend services
- **No search APIs:** Search handled by case-api service
- **No analytics processing:** Analytics handled by timeline-api

Actual Technical Implementation

Core Components

Main Service (`main.go`)

```
// Creates gRPC clients to backend services
caseAPI, err := newCaseAPI() // case-api gRPC client
indexer, err := NewIndexer(caseAPI.CaseClient, caseAPI.TimelineClient, stopper)

// Consumes from Kafka by partition
chans, err := casemconsumer.ConsumeKafkaTopicByPartition[intakecodec.Payload](...)

// Runs parallel workers
for idx, c := range chans {
    go indexer.Run(c, idx) // One goroutine per partition
}
```

Indexer Logic (`indexer.go`)

```
func (i *CaseIndexer) handleRequest(ctx context.Context, r *intakecodec.Payload) error {
    switch req := r.RequestV2.(type) {
```

```

    case *casepb.CreateRequest:
        _, err = i.caseClient.Create(ctx, req) // gRPC call only
    case *casepb.UpdateRequest:
        _, err = i.caseClient.Update(ctx, req) // gRPC call only
    case *timelinepb.CreateCellRequest:
        _, err = i.timelineClient.CreateCell(ctx, req) // gRPC call only
    case *casepb.IntegrationIncrementalSyncRequest:
        _, err = i.caseClient.IntegrationIncrementalSync(ctx, req) // gRPC call only
    }
    return err
}

```

Dependencies

Runtime Dependencies

- **Kafka**: Message broker for `case-intake` topic
- **case-api**: gRPC service for case operations
- **timeline-api**: gRPC service for timeline operations

No Direct Dependencies

- **PostgreSQL**: Connected via case-api/timeline-api only
- **Elasticsearch**: Connected via case-api/timeline-api only
- **External APIs**: All integrations via backend services

Datacenter Deployment Strategy

Multi-Environment Deployment

Deployment Strategy

Figure 4: Deployment Strategy

Deployment Configuration (CORRECTED)

Environment Differences

- **Staging**: `dangerous_everything_parallel` deployment, Monday 5:00 AM
- **Production**: `delta` deployment (DC-by-DC), Monday 8:30 AM
- **Government**: `delta` deployment with FIPS images, Wednesday 8:30 AM

Resource Allocation (All Environments)

```

resources:
  requests:
    cpu: 1 core
    memory: 1Gi
  limits:
    cpu: 1 core
    memory: 2Gi
replicas: 2 # Per datacenter

```

gRPC Configuration

- **Development**: `passthrough://case-api:6481`
- **Production**: TLS-enabled service discovery
- **No database configuration**: All via gRPC

Performance and Observability

Key Metrics Collected

Processing Metrics

- `dd.case_indexer.latency` - Kafka message to gRPC response latency
- `dd.case_indexer.local_retries` - gRPC retry attempts
- `dd.case_indexer.retries` - Pod restart counts for persistent gRPC failures
- `dd.case_indexer.skipped_requests` - Skipped messages by reason
- `dd.case_indexer.error` - gRPC error rates by type and retryability

Business Metrics

- Message **throughput** by Kafka partition
- **gRPC call success/failure rates** by service
- Processing **latency** from Kafka to gRPC completion
- **Duplicate detection** via gRPC status codes

Error Handling Strategy

Three-Tier Error Handling

1. **Local Retries:** Exponential backoff for transient gRPC failures
2. **Pod Restart:** Fatal restart when gRPC services persistently unavailable
3. **Message Skipping:** Skip malformed messages or duplicates (`AlreadyExists`)

Error Classification

```
func isDuplicate(err error) bool {
    if st, ok := status.FromError(err); ok {
        return st.Code() == codes.AlreadyExists // gRPC status code
    }
    return false
}
```

Health Monitoring

Health Checks

- **Liveness:** `/liveness` on port 8080 - service running
- **Readiness:** `/readiness` on port 8080 - ready to process messages
- **No database health checks:** Service is stateless gRPC router

Security and Configuration

Security Features

- **gRPC TLS:** All service-to-service communication encrypted
- **Service Account:** case-management shared RBAC
- **Network Policies:** Cilium-based service isolation
- **No direct database access:** Security handled by backend services

Configuration Management

- **Consul-based:** Dynamic gRPC endpoint configuration
- **Environment-specific:** Dev vs prod gRPC targets
- **No database config:** No connection strings or credentials

Integration Architecture

Upstream (Kafka)

- **case-intake Topic:** 32 partitions
- **Consumer Group:** case-indexer
- **Message Format:** Protobuf via `intakecodec.DecodeIntakeRequest`

Downstream (gRPC Services)

- **case-api:** Case CRUD operations, integration sync
- **timeline-api:** Timeline and analytics operations
- **All Elasticsearch operations:** Handled by these services

No Direct Integrations

- **No Elasticsearch client libraries**
- **No PostgreSQL drivers**
- **No external API clients**
- **No search query processing**

Monitoring and Alerting

Key Monitoring Focus

- **Kafka consumer lag:** Message processing backlog
- **gRPC call latency:** Backend service response times
- **Error rates:** gRPC service availability
- **Pod restart frequency:** Service stability

Dashboard Links

- **Staging:** <https://ddstaging.datadoghq.com/dashboard/rpm-e96-tew>
- **Production:** <https://app.datadoghq.com/dashboard/bnz-bi9-5ar>

Conclusion

The Case-Indexer is a **simple, focused service** that acts as a Kafka-to-gRPC bridge. Its primary value is:

What it Provides

- **Asynchronous processing:** Decouples message publishing from processing
- **Parallel processing:** 32 partition consumers for throughput
- **Reliable delivery:** Retry logic ensures message processing
- **Service isolation:** Separates message consumption from business logic

What it Does NOT Provide

- **Direct database access:** All data operations via gRPC
- **Search capabilities:** All search via case-api
- **Business logic:** Pure message routing
- **Data transformation:** Passes messages as-is

Architecture Benefits

- **Simple:** Easy to understand and maintain
- **Reliable:** Robust error handling and retry logic
- **Scalable:** Partition-based parallel processing
- **Decoupled:** Clean separation between messaging and business logic

The case-indexer serves as a critical but simple component in the case management pipeline - it ensures reliable, asynchronous processing of case events by routing Kafka messages to the appropriate backend services via gRPC.

This corrected report accurately describes the case-indexer as a Kafka-to-gRPC router without direct database or Elasticsearch integration.