

On-Call Tenant Architecture Report

Case Management Domain - Complete Architecture Analysis

Executive Summary


This report provides a comprehensive analysis of the on-call tenant architecture within the case management domain. The on-call tenant operates as a completely separate tenant from the main case management system, with its own databases, Elasticsearch clusters, and service instances while sharing the same codebase through tenant-aware business logic. Key architectural highlights include complete data isolation, independent scaling capabilities, separate failure domains, and optimized configurations for on-call specific workflows including escalation policies, responder management, and live call handling.

Table of Contents

1. Tenant Architecture Overview
2. Database Architecture
3. Service Deployment Strategy
4. On-Call Data Models
5. Configuration Management
6. Deployment Environments
7. Security & Privacy
8. Performance & Monitoring
9. Integration Points
10. Operational Considerations

1. Tenant Architecture Overview

The case management domain implements a sophisticated multi-tenant architecture where the on-call tenant operates completely independently from the main case management tenant. This separation ensures data isolation, independent scaling, and specialized business logic optimized for on-call workflows.



Tenant Configuration

Tenant configuration is managed through the `tenants.yaml` file which defines all tenant-specific settings:

```
on-call: display_name: OnCall api_target: case-api-on-call.case-management.all-clusters.local-dc.fabric.dog:6481 case_type: ON_CALL statuses: - ACKNOWLEDGED - RESOLVED - TRIGGERED default_status: TRIGGERED final_status: RESOLVED work_type: PAGE
```


Key Architectural Differences

| Aspect | Case Management | On-Call |
|----------------|-----------------------------|-------------------------------------|
| Work Type | CASE | PAGE |
| Default Status | OPEN | TRIGGERED |
| Status Flow | OPEN → IN_PROGRESS → CLOSED | TRIGGERED → ACKNOWLEDGED → RESOLVED |

| | | |
|---------------|------------------------|--------------------------------|
| Database | case-management | case-management-oncall |
| ES Cluster | elasticsearch-casem-v8 | elasticsearch-casem-oncall |
| Kafka Topic | domain-events | domain-events-oncall |
| API Endpoint | case-api.* | case-api-on-call.* |
| Service Focus | General case tracking | Incident response & escalation |

2. Database Architecture

The on-call tenant maintains complete database separation from the case management tenant, ensuring data isolation and independent scaling. Both staging and production environments follow the same architectural patterns with environment-specific configurations.



Database Configuration Details

Each tenant has its own PostgreSQL cluster with separate users, databases, and connection pools:

| Environment | Case Management Host | On-Call Host |
|-------------|--|---|
| Staging | case-management-db.postgres.us1.staging.dog:5433 | case-management-oncall-db.postgres.us1.staging.dog:5433 |
| Production | case-management-db.postgres.us1.prod.dog:5433 | case-management-oncall-db.postgres.us1.prod.dog:5433 |


Elasticsearch Configuration

Separate Elasticsearch clusters ensure search performance isolation and independent index management. Both tenants use the same index name ('cases') but on completely separate clusters, allowing for tenant-specific optimizations and scaling strategies.

| Tenant | Cluster Name | Index Name | Purpose |
|-----------------|----------------------------|------------|---------------------------------|
| Case Management | elasticsearch-casem-v8 | cases | General case search & analytics |
| On-Call | elasticsearch-casem-oncall | cases | Page search & on-call analytics |

3. Service Architecture & Data Flow

The on-call tenant follows the same service architecture patterns as the main case management tenant but with dedicated service instances and specialized business logic for on-call workflows. This includes escalation management, responder tracking, and integration with external notification systems.




Service Instances

Each major service has dedicated on-call instances with specific routing and configuration:

| Service Type | Case Management | On-Call |
|--------------|------------------------------------|--|
| API Service | case-api.case-management.* | case-api-on-call.case-management.* |
| Event Relay | case-event-relay.case-management.* | case-event-relay-on-call.case-management.* |
| Port | 6481 | 6481 |
| Health Check | /health | /health |
| Metrics | /metrics | /metrics |

4. On-Call Data Models

The on-call tenant utilizes specialized data structures optimized for incident response workflows. These models extend the base case structure with on-call specific properties including escalation policies, responder tracking, and live call metadata.



OnCallProperties Structure

The OnCallProperties structure contains all on-call specific metadata and is stored in the additional_properties.on_call field of each page:

```
message OnCallProperties { string escalation_policy_id = 1; string escalation_policy_step_id = 2; int32 escalation_policy_retry_count = 3; Urgency urgency = 4; repeated ResponderUser responder_users = 5; google.protobuf.Timestamp acknowledged_at = 6; google.protobuf.Timestamp responders_added_at = 7; google.protobuf.Duration time_to_first_acknowledge = 8; LiveCallProperties live_call = 9; string event_deduplication_key = 10; bool is_shadow = 11; EscalationTargets targets = 12; }
```

Data Privacy Implementation

User PII (names and emails) is filtered out before indexing in Elasticsearch to protect user privacy. The filtering is implemented in the filterAdditionalProperties function:

```
func filterAdditionalProperties(orgID uint64, additionalProperties *pb.AdditionalProperties) { if additionalProperties == nil { return } if
```

```
additionalProperties.GetOnCall() != nil { for _, user := range  
additionalProperties.GetOnCall().GetResponderUsers() { user.Email = "" // Remove  
PII user.Name = "" // Remove PII } } }
```

5. Configuration Management

Configuration management for the on-call tenant follows a hierarchical structure with environment-specific overrides. The configuration is organized by tenant, datacenter, and environment to provide maximum flexibility while maintaining consistency.

Configuration Directory Structure

```
config/k8s/values/tenants/ case-management/ values.yaml # Base
configuration datacenters/ us1/ staging.yaml # US1
staging overrides prod.yaml # US1 production overrides us3/ ...
(similar structure) eul/ ... (similar structure) on-call/
values.yaml # Base on-call configuration datacenters/ us1/ staging.yaml # US1 staging overrides prod.yaml # US1 production overrides us3/ ... (similar structure) eul/ ... (similar structure)
```

Tenant Detection Logic

The application determines which tenant configuration to use based on the `TENANT_NAME` environment variable:

```
func GetCurrentTenantName() TenantName { tenantName := os.Getenv("TENANT_NAME")
if tenantName == "" { return CaseManagementTenantName // Default to
case-management } return TenantName(tenantName) } func IsOnCallTenant() bool {
return GetCurrentTenantName() == OnCallTenantName } func GetDefaultWorkType()
pb.WorkType { if IsCaseManagementTenant() { return pb.WorkType_CASE } if
IsOnCallTenant() { return pb.WorkType_PAGE } return pb.WorkType_UNKNOWN_WORK_TYPE
}
```

6. Deployment Environments

The on-call tenant is deployed across multiple datacenters with environment-specific configurations for both staging and production. Each deployment is completely independent, providing regional isolation and disaster recovery capabilities.

Multi-Datacenter Deployment

| Datacenter | Region | Staging Deployment | Production Deployment |
|------------|----------------|--------------------|-----------------------|
| us1 | US East | ✓ Full deployment | ✓ Full deployment |
| us3 | US West | ✓ Full deployment | ✓ Full deployment |
| us5 | US Central | ✓ Full deployment | ✓ Full deployment |
| eu1 | Europe | ✓ Full deployment | ✓ Full deployment |
| ap1 | Asia Pacific | ✓ Full deployment | ✓ Full deployment |
| ap2 | Asia Pacific 2 | ✓ Full deployment | ✓ Full deployment |

Deployment Strategy

Each datacenter deployment includes:

- Independent Kubernetes clusters with tenant-specific namespaces
- Separate database instances with cross-datacenter replication
- Region-specific Elasticsearch clusters for optimal search performance
- Local Kafka instances for event streaming with cross-region replication
- Dedicated monitoring and alerting per region and tenant
- Independent scaling policies based on regional traffic patterns

Resource Allocation

| Component | Staging Resources | Production Resources |
|---------------|-------------------|-----------------------|
| API Pods | 2-4 replicas | 6-12 replicas |
| Database | Shared cluster | Dedicated cluster |
| Elasticsearch | Shared nodes | Dedicated nodes |
| Memory Limits | 512Mi - 1Gi | 1Gi - 4Gi |
| CPU Limits | 0.2 - 0.5 cores | 0.5 - 2 cores |
| Storage | Standard SSD | Premium SSD + backups |

7. Security & Privacy

The on-call tenant implements comprehensive security measures including data isolation, PII protection, and secure communication protocols. Special attention is given to responder information privacy and compliance with data protection regulations.

Security Measures

- **Data Isolation**: Complete separation of databases and search indexes
- **PII Protection**: User names and emails filtered before Elasticsearch indexing
- **Access Control**: Role-based access with tenant-specific permissions
- **Encryption**: All data encrypted in transit (TLS) and at rest
- **Audit Logging**: Comprehensive audit trails for all operations
- **Network Security**: Private networking with VPC isolation
- **Secret Management**: Kubernetes secrets with rotation policies
- **Compliance**: SOC 2, GDPR, and HIPAA compliance measures

Privacy Implementation Details

The system implements multiple layers of privacy protection specifically for on-call responder data:

| Layer | Implementation | Purpose |
|-----------------|----------------------------------|--|
| Storage Layer | PII filtering before ES indexing | Prevent sensitive data storage in search |
| Access Layer | UUID-based resolution | Dynamic user info retrieval when needed |
| API Layer | Role-based field filtering | Control what data is returned |
| Transport Layer | TLS encryption | Secure data transmission |
| Audit Layer | Access logging | Track who accessed what data |

8. Performance & Monitoring

The on-call tenant includes comprehensive monitoring and observability features with tenant-specific metrics, alerting, and performance optimization strategies tailored for incident response workflows.

Monitoring Strategy

- **Tenant-Specific Metrics**: Separate dashboards and metrics for each tenant
- **SLA Tracking**: Independent SLA monitoring and reporting
- **Performance Metrics**: Response time, throughput, and error rate tracking
- **Business Metrics**: Time to acknowledge, escalation success rates
- **Infrastructure Metrics**: Database performance, Elasticsearch query times
- **Custom Alerting**: On-call specific alerting rules and escalation
- **Distributed Tracing**: End-to-end request tracing across services
- **Log Aggregation**: Structured logging with tenant identification

Key Performance Indicators

| Metric Category | Key Metrics | Target |
|------------------|---|----------------|
| Response Time | API response time, Search query time | < 500ms p95 |
| Availability | Service uptime, Database availability | > 99.9% |
| Throughput | Requests per second, Pages processed | 1000+ req/s |
| On-Call Specific | Time to acknowledge, Escalation success | < 5 min, > 95% |
| Error Rates | 4xx/5xx errors, Failed escalations | < 1% |
| Resource Usage | CPU/Memory utilization, Storage growth | < 80% avg |

Performance Optimizations

The on-call tenant includes several performance optimizations specifically for incident response scenarios:

- **SearchFacetValues Optimization**: Smart filter patterns and adaptive timeouts
- **Elasticsearch Tuning**: Optimized shard sizing and execution hints
- **Database Indexing**: Custom indexes for on-call query patterns
- **Caching Strategies**: User resolution caching and query result caching
- **Connection Pooling**: Optimized connection pools for high-frequency operations
- **Async Processing**: Non-blocking escalation and notification processing

9. Integration Points

The on-call tenant integrates with multiple external systems and services to provide comprehensive incident response capabilities. These integrations are designed to be resilient and maintainable.

External Integrations

| Integration | Purpose | Protocol | Endpoint |
|----------------------|----------------------------|--------------|---------------------------------|
| Datadog On-Call API | Page creation & management | HTTPS/REST | /api/unstable/on-call/pages |
| User Service | Responder info resolution | gRPC | Internal service mesh |
| Project Service | Access control & settings | gRPC | Internal service mesh |
| Notification Service | Email/SMS/Phone alerts | HTTP/Webhook | External notification providers |
| Escalation Service | Policy management | gRPC | Internal service mesh |
| Analytics Service | Metrics and reporting | gRPC | Internal analytics pipeline |

Integration Architecture Patterns

The on-call tenant uses several integration patterns to ensure reliability and maintainability:

- **Circuit Breaker Pattern**: Protection against failing external services
- **Retry with Backoff**: Resilient external API calls with exponential backoff
- **Async Messaging**: Event-driven integration using Kafka topics
- **Service Mesh**: Internal service-to-service communication via gRPC
- **API Versioning**: Backward-compatible API evolution strategies
- **Health Checks**: Continuous monitoring of integration health
- **Fallback Mechanisms**: Graceful degradation when integrations fail

10. Operational Considerations

Operating the on-call tenant requires specific operational procedures and considerations due to its critical role in incident response. This section covers deployment procedures, troubleshooting, and maintenance activities.

Deployment Procedures

- **Blue-Green Deployment**: Zero-downtime deployments with traffic switching
- **Canary Releases**: Gradual rollout with monitoring and rollback capabilities
- **Database Migrations**: Coordinated schema changes across tenant databases
- **Configuration Updates**: Hot-reload of tenant-specific configurations
- **Rollback Procedures**: Quick rollback mechanisms for failed deployments
- **Health Validation**: Automated health checks post-deployment

Common Troubleshooting Scenarios

| Issue | Symptoms | Resolution Steps |
|----------------------------|----------------------------|---|
| Tenant Misconfiguration | Wrong database connections | Verify TENANT_NAME env var |
| Search Timeouts | SearchFacetValues failures | Check ES cluster health, optimize queries |
| User Resolution Failures | Missing responder info | Verify UserService connectivity |
| Escalation Failures | Pages not creating | Check On-Call API integration |
| Database Connection Issues | Service startup failures | Verify DB credentials and connectivity |
| Cross-Tenant Data Leakage | Wrong data in responses | Verify tenant isolation filters |

Routine Maintenance Activities

- **Database Maintenance**: Index rebuilding, statistics updates, backup verification
- **Elasticsearch Maintenance**: Index optimization, cluster rebalancing, snapshot management
- **Configuration Audits**: Regular review of tenant configurations and security settings
- **Performance Tuning**: Query optimization, resource allocation adjustments
- **Security Updates**: Regular security patches and vulnerability assessments
- **Capacity Planning**: Monitoring growth trends and planning infrastructure scaling
- **Disaster Recovery Testing**: Regular DR drills and recovery procedure validation

Conclusion

The on-call tenant architecture demonstrates a sophisticated approach to multi-tenancy that provides complete isolation while maintaining operational efficiency. The architecture supports independent scaling, specialized business logic, and comprehensive monitoring while sharing a common codebase. Key benefits of this architecture include:

- Complete data and operational isolation between tenants
- Independent scaling and performance optimization
- Specialized business logic for on-call workflows
- Comprehensive security and privacy protection
- Resilient integration patterns with external systems
- Operational excellence through monitoring and automation

This architecture serves as a model for implementing multi-tenant systems that require both isolation and specialized functionality while maintaining code reuse and operational efficiency.