

Case Management Search Architecture

Technical Analysis Report

Comprehensive Analysis of Search Features,
Elasticsearch Integration, and Analytics Module
in the Case Management Domain

Generated on: 2025-09-19 10:38:27

Datadog Case Management Team

Key Statistics

- 393+ Elasticsearch fields
- 7 specialized query handlers
- Nested custom attributes support
- Multi-environment deployment (US1, EU1, AP1+)

Executive Summary

The Case Management domain implements a sophisticated search architecture built on Elasticsearch, providing comprehensive search capabilities for case data across multiple Datadog environments.

KEY FINDINGS:

□ SEARCH ARCHITECTURE

- Multi-layered architecture: API → Handlers → Parser → Elasticsearch
- 7 specialized query handlers for different search use cases
- ANTLR-based query parsing with custom grammar
- Project-based security filtering integrated at query level

□ ELASTICSEARCH INTEGRATION

- Primary index: "cases" with 393+ mapped fields
- Custom attributes stored as nested objects (not flattened)
- Dynamic runtime field support for complex calculations
- Multi-environment deployment across US, EU, and AP regions

□ ANALYTICS MODULE

- Real-time time-series aggregations with configurable intervals
- Support for 20+ groupBy fields plus custom attributes
- Complex nested aggregations for custom attribute analysis
- Metrics: count, sum, avg, max, min, percentiles (pc50, pc95, pc99)

□ ADVANCED FEATURES

- Faceted search with autocomplete
- Custom attribute querying via nested Elasticsearch structures
- User/project/case-type enrichment with UUID resolution
- Pagination, sorting, and filtering capabilities

□ PERFORMANCE OPTIMIZATIONS

- Query builder with project restriction filtering
- Aggregation bucket limits (max 1000 buckets)
- Caching and timeout configurations
- Multiple search indexer applications for real-time updates

The system demonstrates enterprise-grade search capabilities with strong separation of concerns, robust error handling, and comprehensive logging for debugging and monitoring.

Search Architecture Overview

COMPONENT BREAKDOWN:

1. CLIENT LAYER

- Web UI search interfaces
- API clients (REST/gRPC)
- External integrations

2. API GATEWAY

- Case Rapid API (/search endpoints)
- Request validation and routing
- Authentication and authorization

3. QUERY HANDLERS (7 specialized handlers)

- SearchHandler: Basic search with pagination/sorting
- AnalyticHandler: Time-series aggregations and metrics
- SearchFacetValuesHandler: Faceted search with partial matching
- SearchAutocompleteHandler: Search suggestions
- GetFacetsHandler: Available facets retrieval
- GetAttributeValuesHandler: Attribute value enumeration
- GetAttributesKeyHandler: Attribute key discovery

4. QUERY PROCESSING

- ANTLR parser with custom grammar (grammar.peg)
- ESQueryBuilder converts parsed queries to Elasticsearch DSL
- Support for complex nested queries and custom attributes
- Project-based access control filtering

5. ELASTICSEARCH CLUSTER

- Primary index: "cases" (393+ fields)
- Nested custom_attributes structure
- Analytics fields with date_range spans
- Multi-environment deployment

6. ENRICHMENT SERVICES

- UserService (OUI integration) - UUID to email resolution
- ProjectService - Project ID to name mapping
- CaseTypeService - Case type enrichment
- Real-time data enrichment post-query

Elasticsearch Index Structure

INDEX: "cases" (393+ Fields)

CORE FIELDS:

- internal_id, case_id, public_id (keyword) - Unique identifiers
- assignee_id, created_by, project_id, org_id (keyword/long) - References
- status, priority, type_id (long) - Enumerated values
- created_at, modified_at, closed_at (date) - Timestamps
- title, description (text) - Full-text searchable content

CUSTOM ATTRIBUTES (NESTED):

- Type: "nested" (enables complex nested queries)
- Structure:
 - key (keyword) - Attribute name
 - value_text (text) - String values
 - value_number (text + as_double field) - Numeric values
- Query Pattern: nested → filter → terms → reverse_nested

ANALYTICS FIELDS:

- analytics.status_* (object) - Status duration tracking
- spans (date_range) - Time period definitions
- total (long) - Calculated duration metrics
- Supports time-series aggregations

ADDITIONAL PROPERTIES:

- event_management.* - Incident management data
- on_call.* - Responder and escalation information
- change_request.* - Change request metadata
- campaign.* - Campaign tracking data
- Flattened structure for direct querying

INTEGRATION FIELDS:

- jira_issue.* - Jira integration metadata
- servicenow_ticket.* - ServiceNow synchronization
- insights.* - ML/AI insight references
- notification_handles.* - Alert routing information

PERFORMANCE CONSIDERATIONS:

- Dynamic mapping: false (explicit schema control)
- Text fields: analyzed for full-text search
- Keyword fields: not_analyzed for exact matching
- Date fields: optimized for range queries
- Nested fields: isolated document storage for complex queries

Analytics Module Integration

ANALYTICS HANDLER (analytic_handler.go:1520 lines)

SUPPORTED METRICS:

- count - Document counting
- sum, avg, max, min - Statistical aggregations
- pc50, pc95, pc99 - Percentile calculations
- Custom metric definitions via metrics framework

GROUP BY FIELDS (20+ supported):

- Standard: assignee, creator, status, priority, project, service, team
- Time-based: created_at, modified_at, closed_at
- Custom: custom_attributes.* (any custom attribute key)
- Advanced: change_request.risk, event_management.event_aggregation_keys

TIME BUCKETING:

- Configurable intervals: 1s to 1year
- Default intervals: [1s, 2s, 5s, 10s, 1m, 2m, 5m, 10m, 1h, 1d, 7d, 30d, 365d]
- Max bucket limit: 1000 buckets (configurable)
- Automatic interval adjustment to stay under limit

AGGREGATION STRUCTURE:

1. Group By Aggregation (if specified)
 - Terms/MultiTerms for standard fields
 - Nested aggregation for custom attributes
2. Time Bucket Aggregation
 - DateRange aggregation with calculated intervals
3. Metric Aggregation (if specified)
 - Applied within each time bucket

CUSTOM ATTRIBUTE SUPPORT:

- Complex nested aggregation chain:
nested → filter(key match) → terms(values) → reverse_nested → date_range → metric
- Supports both text and numeric custom attribute values
- text_values: uses .keyword field for exact matching
- number_values: uses .as_double field for numeric operations

ENRICHMENT PROCESS:

- UUID to human-readable name conversion
- User resolution via OUI service (UserService)
- Project name resolution (ProjectService)
- Case type name resolution (CaseTypeService)
- Applied post-aggregation for performance

RESPONSE FORMAT:

- MetricsBuffer with LZ4 compression
- Time-series data optimized for frontend consumption
- Groups, epochs, and values arrays
- Supports Datadog's metrics infrastructure

Query Processing Flow

QUERY PROCESSING PIPELINE:

1. INPUT PARSING (ANTLR Parser)

- Grammar-based parsing (grammar.peg)
- Supports complex boolean logic: AND, OR, NOT
- Field-specific queries: status:open, priority:high
- Range queries: created_at:[2023-01-01 TO 2023-12-31]
- Custom attribute queries: custom_attributes.environment:prod

2. QUERY BUILDING (ESQueryBuilder)

- Converts parsed tree to Elasticsearch Query DSL
- Handles different query types:
 - Match queries for full-text search
 - Term queries for exact matching
 - Range queries for numeric/date ranges
 - Nested queries for custom attributes
 - Boolean queries for complex logic

3. SECURITY FILTERING

- Project-based access control
- Automatic org_id filtering
- Restricted project ID filtering via ProjectService
- Work type filtering (CASE, etc.)

4. ELASTICSEARCH EXECUTION

- Query optimization and caching
- Timeout configuration (5s for analytics)
- Index selection based on configuration
- Result aggregation and processing

5. RESULT PROCESSING

- Hit extraction from Elasticsearch response
- Document deserialization (JSON → protobuf)
- Pagination and sorting application
- Total count and page count calculation

6. ENRICHMENT

- UUID resolution to human names
- Project name enrichment
- User email/handle resolution
- Case type name mapping

7. RESPONSE FORMATTING

- Protobuf serialization
- Metadata addition (totals, pagination info)
- Error handling and logging

SUPPORTED QUERY PATTERNS:

- Simple: "status:open"
- Boolean: "status:open AND priority:high"
- Range: "created_at:[now-7d TO now]"
- Custom attributes: "custom_attributes.environment:production"
- Complex: "(status:open OR status:in_progress) AND assignee:user123"

Deployment Architecture

MULTI-ENVIRONMENT DEPLOYMENT:

PRODUCTION ENVIRONMENTS:

- US1 (us-east-1) - Primary US production
- US3 (us-west-2) - Secondary US production
- US5 (us-west-1) - Tertiary US production
- EU1 (eu-west-1) - European production
- AP1 (ap-southeast-1) - Asia Pacific production
- AP2 (ap-northeast-1) - Secondary AP production

STAGING ENVIRONMENT:

- US1-STAGING - Full staging environment with production data subset

SEARCH INDEXER APPLICATIONS:

Multiple case-event-handler deployments per environment:

- case-event-handler-search-indexer-es-v8-us1.yaml
- case-event-handler-search-indexer-es-v8-eu1.yaml
- case-event-handler-search-indexer-es-v8-ap1.yaml
- case-event-handler-analytic-*.yaml (analytics configuration)

ELASTICSEARCH CLUSTERS:

- Environment-specific ES clusters
- Index per environment isolation
- Cluster-level security and access control
- Regional data sovereignty compliance

CONFIGURATION MANAGEMENT:

- Consul-based configuration (consulconfig)
- Environment-specific overrides
- Feature flags and runtime parameters
- Monitoring and alerting integration

KEY APPLICATIONS:

1. CASE-RAPID-API

- REST API endpoints for search
- gRPC service interfaces
- Authentication and rate limiting
- Request/response transformation

2. CASE-EVENT-HANDLER

- Real-time index updates
- Event sourcing integration
- Search index synchronization
- Analytics data computation


3. SEARCH INDEXER

- Elasticsearch document management
- Index lifecycle management
- Mapping updates and migrations
- Performance optimization

MONITORING & OBSERVABILITY:

- Structured logging with contextual information
- Metrics collection and dashboards
- Error tracking and alerting
- Performance monitoring and profiling

Search Architecture Diagram



Elasticsearch Index Mapping

Elasticsearch Index Structure: "cases"

Core Case Fields

- internal_id, case_id, public_id (keyword)
- assignee_id, created_by (keyword)
- project_id, org_id (keyword/long)
- status, priority, type_id (long)
- created_at, modified_at (date)
- title, description (text)

Custom Attributes (Nested)

- Type: "nested" (not flattened object)
- key (keyword) - attribute name
- value_text (text) - string values
- value_number (text + as_double) - numeric
- Enables complex nested queries

Analytics Fields

- analytics.status_* (object)
- spans (date_range) - time periods
- total (long) - duration metrics
- Supports time-based aggregations
- Pre-computed metrics storage

Additional Properties

- event_management.* - incident data
- on_call.* - responder information
- change_request.* - CR metadata
- campaign.* - campaign data
- Flattened structure for easy querying

Standard Attributes

- attributes (flattened) - key-value pairs
- service, team, version, datacenter
- rule_id, follow_up_incident_public_id
- Used for faceted search and filtering
- Direct field access for performance

Integration Fields

- jira_issue.* - Jira integration
- servicenow_ticket.* - ServiceNow
- insights.* - ML insights
- notification_handles.* - alerting
- External system synchronization

Index Mapping Statistics

- Total Properties: 393+ fields
- Dynamic Mapping: false (explicit schema)
- Nested Objects: custom_attributes
- Flattened Objects: attributes, additional_properties


- Text Fields: title, description (full-text search)
- Keyword Fields: IDs, status enums (exact match)
- Date Fields: timestamps with time-based queries
- Long Fields: counters, enums, numeric identifiers

Query Processing Flow

Search Query Processing Flow



Analytics Aggregation Structure



Custom Attribute Aggregation Details

Structure: nested → filter → terms → reverse_nested → date_range → metric

1. Nested: Enter custom_attributes array
2. Filter: Match specific attribute key (e.g., "environment")
3. Terms: Group by text_values.keyword or number_values.as_double
4. Reverse_nested: Exit to parent document for time calculations
5. Date_range: Create time buckets
6. Metric: Apply aggregation method (count, avg, etc.)

Conclusions & Recommendations

STRENGTHS:

□ ROBUST ARCHITECTURE

- Clean separation of concerns with specialized handlers
- Comprehensive error handling and logging
- Scalable multi-environment deployment
- Strong typing with protobuf integration

□ ADVANCED SEARCH CAPABILITIES

- Complex query parsing with ANTLR grammar
- Nested custom attribute support
- Real-time analytics with flexible aggregations
- Comprehensive faceted search functionality

□ PERFORMANCE OPTIMIZATIONS

- Project-based security filtering
- Query builder optimizations
- Configurable timeouts and limits
- LZ4 compression for analytics responses

RECOMMENDATIONS FOR IMPROVEMENT:

□ PERFORMANCE ENHANCEMENTS

- Implement query result caching for frequently accessed data
- Add query performance monitoring and slow query alerts
- Consider read replicas for analytics-heavy workloads
- Optimize aggregation queries for large datasets

□ SCALABILITY IMPROVEMENTS

- Implement horizontal scaling for search handlers
- Add connection pooling for Elasticsearch clients
- Consider implementing circuit breakers for external services
- Add request queuing and rate limiting

□ MONITORING & OBSERVABILITY

- Enhanced metrics collection for query performance
- Distributed tracing for complex query flows
- Custom dashboards for search analytics
- Automated alerting for search failures

□ DOCUMENTATION & TOOLING

- Interactive query builder UI for testing
- Comprehensive API documentation with examples
- Performance testing framework
- Schema migration tools for Elasticsearch mappings

FUTURE CONSIDERATIONS:

- Machine learning integration for search ranking
- Support for vector/semantic search capabilities
- Advanced caching strategies (Redis integration)
- Multi-tenancy improvements for better isolation